

In [1]:

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```

In []:

In [5]:

```
#Loading the dataset to a pandas Dataframe
sonar_data = pd.read_csv('Copy of sonar data.csv', header=None)
```

In [6]:

```
sonar_data.head()
```

Out[6]:

	0	1	2	3	4	5	6	7	8	9	...	51	
0	0.0200	0.0371	0.0428	0.0207	0.0954	0.0986	0.1539	0.1601	0.3109	0.2111	...	0.0027	0.
1	0.0453	0.0523	0.0843	0.0689	0.1183	0.2583	0.2156	0.3481	0.3337	0.2872	...	0.0084	0.
2	0.0262	0.0582	0.1099	0.1083	0.0974	0.2280	0.2431	0.3771	0.5598	0.6194	...	0.0232	0.
3	0.0100	0.0171	0.0623	0.0205	0.0205	0.0368	0.1098	0.1276	0.0598	0.1264	...	0.0121	0.
4	0.0762	0.0666	0.0481	0.0394	0.0590	0.0649	0.1209	0.2467	0.3564	0.4459	...	0.0031	0.

5 rows × 61 columns

In [7]:

```
# number of rows and columns
sonar_data.shape
```

Out[7]:

(208, 61)

In [8]:

```
sonar_data.describe() #describe --> statistical measures of the data
```

Out[8]:

	0	1	2	3	4	5	6	
count	208.000000	208.000000	208.000000	208.000000	208.000000	208.000000	208.000000	20
mean	0.029164	0.038437	0.043832	0.053892	0.075202	0.104570	0.121747	
std	0.022991	0.032960	0.038428	0.046528	0.055552	0.059105	0.061788	
min	0.001500	0.000600	0.001500	0.005800	0.006700	0.010200	0.003300	
25%	0.013350	0.016450	0.018950	0.024375	0.038050	0.067025	0.080900	
50%	0.022800	0.030800	0.034300	0.044050	0.062500	0.092150	0.106950	
75%	0.035550	0.047950	0.057950	0.064500	0.100275	0.134125	0.154000	
max	0.137100	0.233900	0.305900	0.426400	0.401000	0.382300	0.372900	

8 rows × 60 columns

In [9]:

```
sonar_data[60].value_counts()
```

Out[9]:

```
M    111
R     97
Name: 60, dtype: int64
```

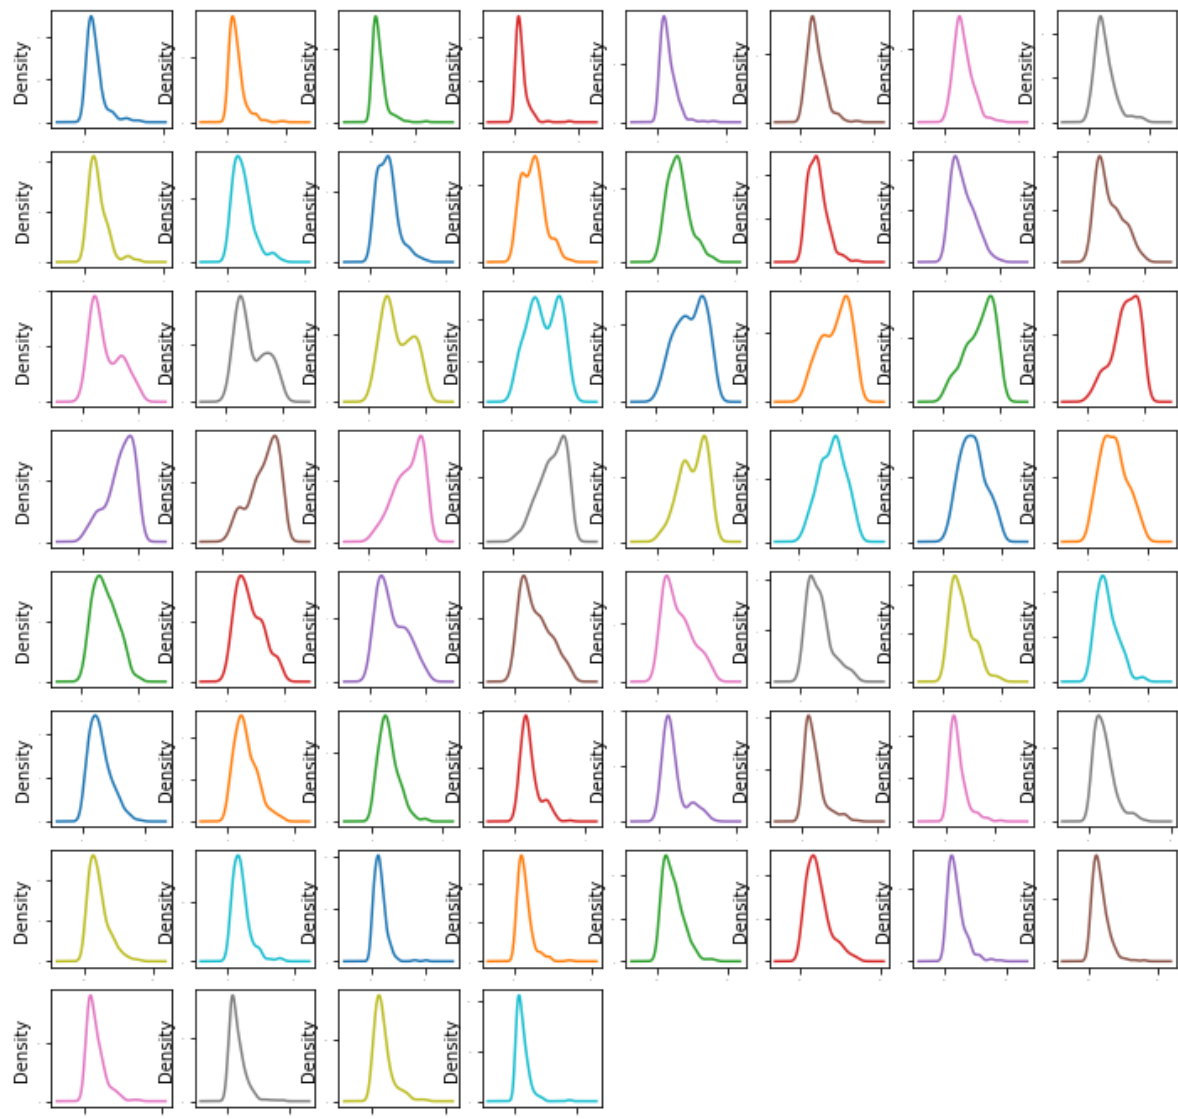
In [28]:

```
from matplotlib import pyplot
sonar_data.hist(sharex=False, sharey=False, xlabelsize=1, ylabelsize=1, figsize=(12,12))
pyplot.show()
```



In [29]:

```
sonar_data.plot(kind='density',subplots=True,layout=(8,8),sharex=False,legend=False,fontsize=10,figsize=(16,16))
```



In [10]:

```
sonar_data.groupby(60).mean()
```

Out[10]:

	0	1	2	3	4	5	6	7	8
60									
M	0.034989	0.045544	0.050720	0.064768	0.086715	0.111864	0.128359	0.149832	0.213492
R	0.022498	0.030303	0.035951	0.041447	0.062028	0.096224	0.114180	0.117596	0.137392

2 rows × 60 columns

In [11]:

```
# separating data and Labels  
X = sonar_data.drop(columns=60, axis=1)  
Y = sonar_data[60]
```

In [12]:

```
print(X)
```

\	0	1	2	3	4	5	6	7	8
0	0.0200	0.0371	0.0428	0.0207	0.0954	0.0986	0.1539	0.1601	0.3109
1	0.0453	0.0523	0.0843	0.0689	0.1183	0.2583	0.2156	0.3481	0.3337
2	0.0262	0.0582	0.1099	0.1083	0.0974	0.2280	0.2431	0.3771	0.5598
3	0.0100	0.0171	0.0623	0.0205	0.0205	0.0368	0.1098	0.1276	0.0598
4	0.0762	0.0666	0.0481	0.0394	0.0590	0.0649	0.1209	0.2467	0.3564
..
203	0.0187	0.0346	0.0168	0.0177	0.0393	0.1630	0.2028	0.1694	0.2328
204	0.0323	0.0101	0.0298	0.0564	0.0760	0.0958	0.0990	0.1018	0.1030
205	0.0522	0.0437	0.0180	0.0292	0.0351	0.1171	0.1257	0.1178	0.1258
206	0.0303	0.0353	0.0490	0.0608	0.0167	0.1354	0.1465	0.1123	0.1945
207	0.0260	0.0363	0.0136	0.0272	0.0214	0.0338	0.0655	0.1400	0.1843
	9	...	50	51	52	53	54	55	56
0	0.2111	...	0.0232	0.0027	0.0065	0.0159	0.0072	0.0167	0.0180
1	0.2872	...	0.0125	0.0084	0.0089	0.0048	0.0094	0.0191	0.0140
2	0.6194	...	0.0033	0.0232	0.0166	0.0095	0.0180	0.0244	0.0316
3	0.1264	...	0.0241	0.0121	0.0036	0.0150	0.0085	0.0073	0.0050
4	0.4459	...	0.0156	0.0031	0.0054	0.0105	0.0110	0.0015	0.0072
..
203	0.2684	...	0.0203	0.0116	0.0098	0.0199	0.0033	0.0101	0.0065
204	0.2154	...	0.0051	0.0061	0.0093	0.0135	0.0063	0.0063	0.0034
205	0.2529	...	0.0155	0.0160	0.0029	0.0051	0.0062	0.0089	0.0140
206	0.2354	...	0.0042	0.0086	0.0046	0.0126	0.0036	0.0035	0.0034
207	0.2354	...	0.0181	0.0146	0.0129	0.0047	0.0039	0.0061	0.0040
	57	58	59						
0	0.0084	0.0090	0.0032						
1	0.0049	0.0052	0.0044						
2	0.0164	0.0095	0.0078						
3	0.0044	0.0040	0.0117						
4	0.0048	0.0107	0.0094						
..						
203	0.0115	0.0193	0.0157						
204	0.0032	0.0062	0.0067						
205	0.0138	0.0077	0.0031						
206	0.0079	0.0036	0.0048						
207	0.0036	0.0061	0.0115						

[208 rows x 60 columns]



In [13]:

```
print(Y)
```

```
0      R
1      R
2      R
3      R
4      R
..
203    M
204    M
205    M
206    M
207    M
```

Name: 60, Length: 208, dtype: object

In [14]:

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.1, stratify=Y, rand
```

In [15]:

```
print(X.shape, X_train.shape, X_test.shape)
```

```
(208, 60) (187, 60) (21, 60)
```

In [16]:

```
print(X_train)
```

	0	1	2	3	4	5	6	7	8
\									
115	0.0414	0.0436	0.0447	0.0844	0.0419	0.1215	0.2002	0.1516	0.0818
38	0.0123	0.0022	0.0196	0.0206	0.0180	0.0492	0.0033	0.0398	0.0791
56	0.0152	0.0102	0.0113	0.0263	0.0097	0.0391	0.0857	0.0915	0.0949
123	0.0270	0.0163	0.0341	0.0247	0.0822	0.1256	0.1323	0.1584	0.2017
18	0.0270	0.0092	0.0145	0.0278	0.0412	0.0757	0.1026	0.1138	0.0794
..
140	0.0412	0.1135	0.0518	0.0232	0.0646	0.1124	0.1787	0.2407	0.2682
5	0.0286	0.0453	0.0277	0.0174	0.0384	0.0990	0.1201	0.1833	0.2105
154	0.0117	0.0069	0.0279	0.0583	0.0915	0.1267	0.1577	0.1927	0.2361
131	0.1150	0.1163	0.0866	0.0358	0.0232	0.1267	0.2417	0.2661	0.4346
203	0.0187	0.0346	0.0168	0.0177	0.0393	0.1630	0.2028	0.1694	0.2328
	9	...	50	51	52	53	54	55	56
115	0.1975	...	0.0222	0.0045	0.0136	0.0113	0.0053	0.0165	0.0141
38	0.0475	...	0.0149	0.0125	0.0134	0.0026	0.0038	0.0018	0.0113
56	0.1504	...	0.0048	0.0049	0.0041	0.0036	0.0013	0.0046	0.0037
123	0.2122	...	0.0197	0.0189	0.0204	0.0085	0.0043	0.0092	0.0138
18	0.1520	...	0.0045	0.0084	0.0010	0.0018	0.0068	0.0039	0.0120
..
140	0.2058	...	0.0798	0.0376	0.0143	0.0272	0.0127	0.0166	0.0095
5	0.3039	...	0.0104	0.0045	0.0014	0.0038	0.0013	0.0089	0.0057
154	0.2169	...	0.0039	0.0053	0.0029	0.0020	0.0013	0.0029	0.0020
131	0.5378	...	0.0228	0.0099	0.0065	0.0085	0.0166	0.0110	0.0190
203	0.2684	...	0.0203	0.0116	0.0098	0.0199	0.0033	0.0101	0.0065
	57	58	59						
115	0.0077	0.0246	0.0198						
38	0.0058	0.0047	0.0071						
56	0.0011	0.0034	0.0033						
123	0.0094	0.0105	0.0093						
18	0.0132	0.0070	0.0088						
..						
140	0.0225	0.0098	0.0085						
5	0.0027	0.0051	0.0062						
154	0.0062	0.0026	0.0052						
131	0.0141	0.0068	0.0086						
203	0.0115	0.0193	0.0157						

[187 rows x 60 columns]

In [17]:

```
print(Y_train)
```

```
115    M
38     R
56     R
123    M
18     R
..
140    M
5      R
154    M
131    M
203    M
Name: 60, Length: 187, dtype: object
```

In [18]:

```
model = LogisticRegression()
```

In [19]:

```
#training the Logistic Regression model with training data
model.fit(X_train, Y_train)
```

Out[19]:

```
LogisticRegression()
```

In [20]:

```
#accuracy on training data
X_train_prediction = model.predict(X_train)
training_data_accuracy = accuracy_score(X_train_prediction, Y_train)
```

In [21]:

```
print('Accuracy on training data : ', training_data_accuracy)
```

```
Accuracy on training data :  0.8342245989304813
```

In [22]:

```
#accuracy on test data
X_test_prediction = model.predict(X_test)
test_data_accuracy = accuracy_score(X_test_prediction, Y_test)
```

In [23]:

```
print('Accuracy on test data : ', test_data_accuracy)
```

```
Accuracy on test data :  0.7619047619047619
```

In [24]:

```
input_data = (0.0307,0.0523,0.0653,0.0521,0.0611,0.0577,0.0665,0.0664,0.1460,0.2792,0.3877,  
  
# changing the input_data to a numpy array  
input_data_as_numpy_array = np.asarray(input_data)  
  
# reshape the np array as we are predicting for one instance  
input_data_resaped = input_data_as_numpy_array.reshape(1,-1)  
  
prediction = model.predict(input_data_resaped)  
print(prediction)  
  
if (prediction[0]=='R'):  
    print('The object is a Rock')  
else:  
    print('The object is a mine')
```

['M']
The object is a mine

In []: